

# Employee Turnover

January 3, 2025

```
[10]: import pandas as pd
import numpy as np
import scipy.stats
import seaborn as sns
from sklearn.metrics.pairwise import cosine_similarity
import matplotlib.pyplot as plt
import statistics
import operator
from sklearn.metrics import silhouette_score
from imblearn.pipeline import Pipeline
from sklearn.model_selection import StratifiedKFold
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE
from sklearn.cluster import KMeans
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression # Logistic Regression
from sklearn.ensemble import RandomForestClassifier # Random Forest
from sklearn.ensemble import GradientBoostingClassifier # Gradient Boosting
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import roc_curve
from sklearn.metrics import auc
from sklearn.metrics import classification_report
```

0.0.1 1.

0.0.2 Data Quality Check

```
[11]: df = pd.read_csv('HR_comma_sep.csv')
df.head()
```

```
[11]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	

	time_spend_company	Work_accident	left	promotion_last_5years	sales	\
--	--------------------	---------------	------	-----------------------	-------	---

0	3	0	1	0	sales
1	6	0	1	0	sales
2	4	0	1	0	sales
3	5	0	1	0	sales
4	3	0	1	0	sales

```

salary
0    low
1  medium
2  medium
3    low
4    low

```

```
[12]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   satisfaction_level     14999 non-null float64
1   last_evaluation       14999 non-null float64
2   number_project        14999 non-null int64
3   average_monthly_hours 14999 non-null int64
4   time_spend_company    14999 non-null int64
5   Work_accident         14999 non-null int64
6   left                 14999 non-null int64
7   promotion_last_5years 14999 non-null int64
8   sales                 14999 non-null object
9   salary               14999 non-null object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB

```

### Check for missing values

```
[13]: df.isnull().sum()
```

```

[13]: satisfaction_level    0
      last_evaluation      0
      number_project       0
      average_monthly_hours 0
      time_spend_company    0
      Work_accident        0
      left                 0
      promotion_last_5years 0
      sales                0
      salary               0
      dtype: int64

```

### Drop duplicated records

```
[14]: df.duplicated().sum()  
df.shape
```

```
[14]: (14999, 10)
```

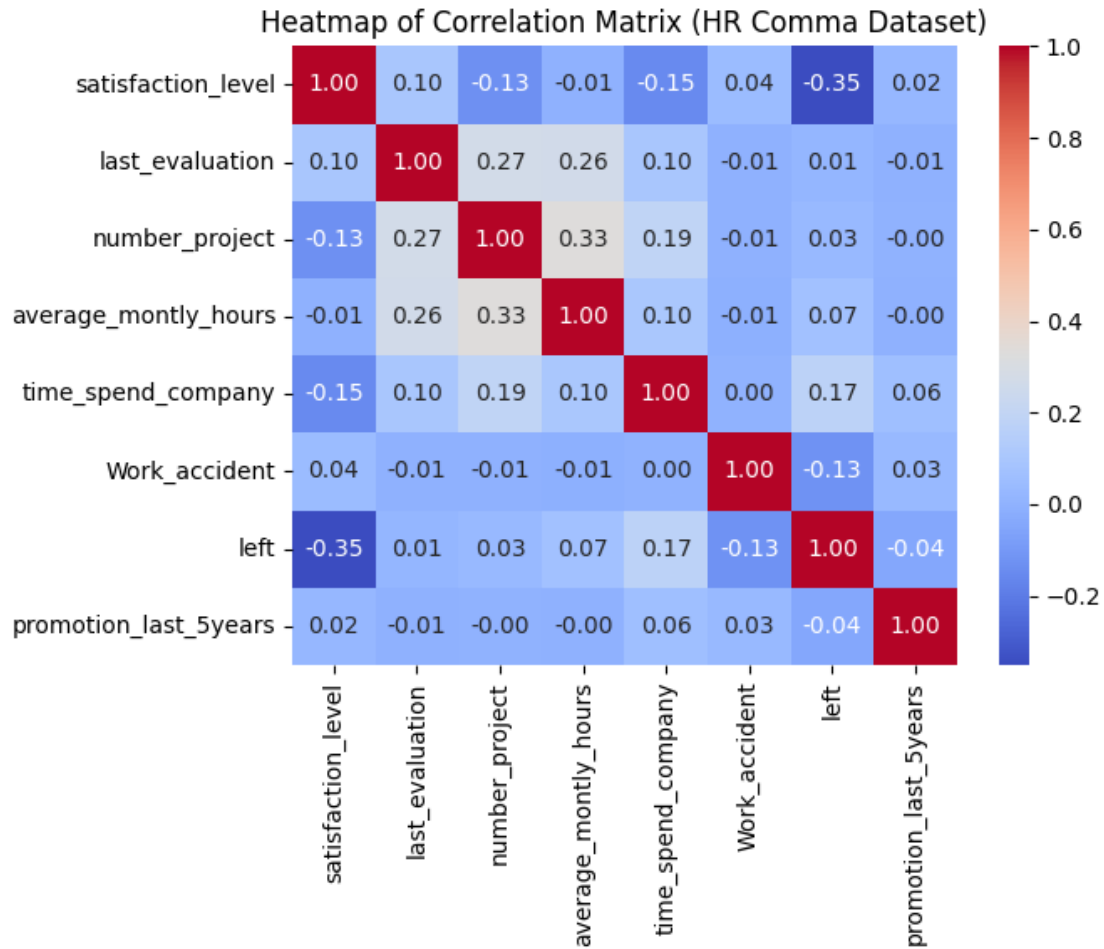
```
[15]: #drop duplicate records  
df = df.drop_duplicates()  
df.shape
```

```
[15]: (11991, 10)
```

### 0.0.3 2.

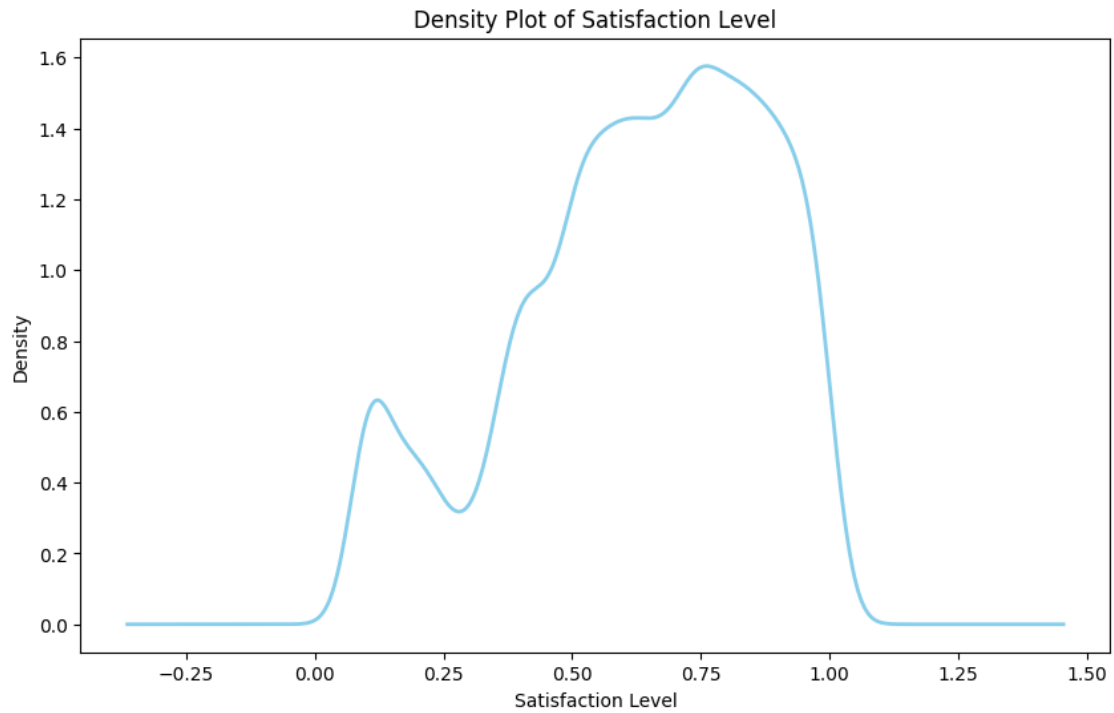
#### Heatmap of correlation between all numerical features in the data

```
[16]: # heatmap of correlation of all numeric features.  
correlation_matrix = df.select_dtypes(include='number').corr()  
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")  
plt.title('Heatmap of Correlation Matrix (HR Comma Dataset)')  
plt.show()
```



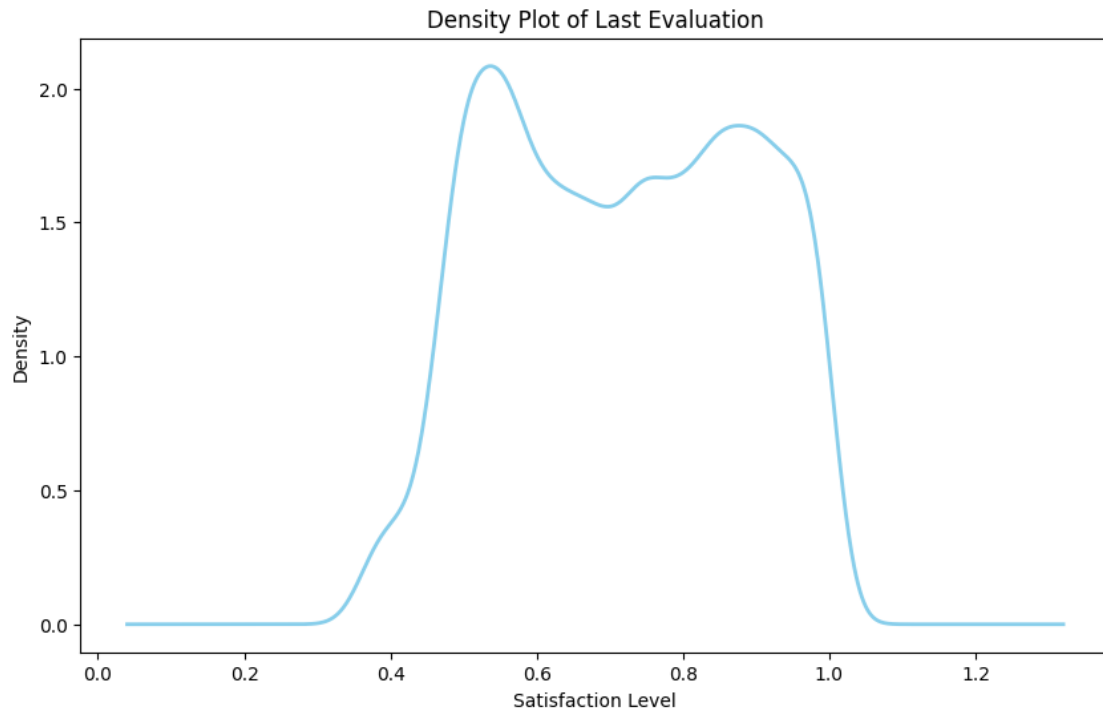
### Employee Satisfaction Distribution Plot

```
[17]: # density plot best shows the relationships between what we are looking for
plt.figure(figsize=(10, 6))
df['satisfaction_level'].plot(kind='density', color='skyblue', linewidth=2)
plt.title('Density Plot of Satisfaction Level')
plt.xlabel('Satisfaction Level')
plt.ylabel('Density')
plt.show()
```



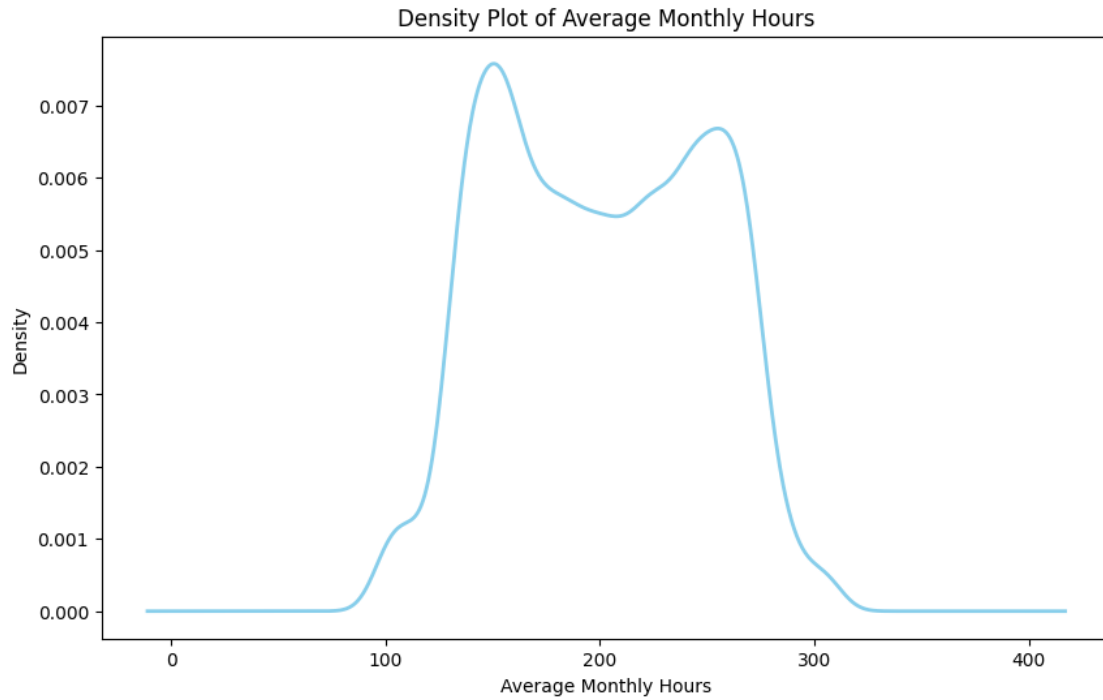
### Employee Evaluation Distribution Plot

```
[18]: plt.figure(figsize=(10, 6))
df['last_evaluation'].plot(kind='density', color='skyblue', linewidth=2)
plt.title('Density Plot of Last Evaluation')
plt.xlabel('Satisfaction Level')
plt.ylabel('Density')
plt.show()
```



### Employee Average Monthly Hours Distribution Plot


```
[19]: plt.figure(figsize=(10, 6))
df['average_monthly_hours'].plot(kind='density', color='skyblue', linewidth=2)
plt.title('Density Plot of Average Monthly Hours')
plt.xlabel('Average Monthly Hours')
plt.ylabel('Density')
plt.show()
```



Bar plot of the employee project count of both employees who left and stayed in the organization

```
[20]: # Bar Plot with count of employees for each number_project differentiated by
      ↪ left
      sns.countplot(x='number_project', data=df, hue='left')
      plt.title('Bar Plot: Count of Employees by Number of Projects (Differentiated
      ↪ by Left)')
      plt.ylabel('Count of Employees')
      plt.xlabel('Number of Projects')
      plt.show()
```

Bar Plot: Count of Employees by Number of Projects (Differentiated by Left)



**Underwork and Overwork Cause Turnover:** Employees with 2 projects (underworked) or 5+ projects (overworked) are more likely to leave. Ensuring an optimal workload of 3-4 projects can reduce turnover.

### 0.0.4 3.

### 0.0.5 Perform clustering of employees who left based on their satisfaction and evaluation


Finding optimal K for K-means clustering of employees who left the company based on their satisfaction and evaluation

```
[21]: satisfaction_evaluation_left = df[df['left'] == 1][['satisfaction_level',
    ↪ 'last_evaluation']]
```

```
[22]: wcss = []
    for i in range(1, 11):
        model = KMeans(n_clusters = i, n_init=10, init = 'k-means++', random_state=
    ↪ 42)
        model.fit(satisfaction_evaluation_left)
        wcss.append(model.inertia_)
    plt.plot(range(1, 11), wcss)
    plt.title('The Elbow Method')
```




```
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()
```



```
[23]: #since using wcss score was tricky, we can use silhouette score to find the
      ↪ optimal number of clusters :
      # Silhouette Score method
      silhouette_scores = []
      for i in range(2, 11):
          model = KMeans(n_clusters=i, n_init=10, init='k-means++', random_state=42)
          model.fit(satisfaction_evaluation_left)
          score = silhouette_score(satisfaction_evaluation_left, model.labels_)
          silhouette_scores.append(score)

      plt.plot(range(2, 11), silhouette_scores)
      plt.title('Silhouette Score Method')
      plt.xlabel('Number of clusters')
      plt.ylabel('Silhouette Score')
      plt.show()
```

```
# we see that three is the optimal number of clusters, so we go ahead and make
↳ three clusters
```



```
[24]: optimal_clusters = range(2, 11)[silhouette_scores.index(max(silhouette_scores))]
      optimal_clusters
```

```
[24]: 3
```

### 0.0.6 Optimal number of K = 3

```
[25]: model = KMeans(n_clusters=optimal_clusters, n_init=10, init='k-means++',
↳ random_state=42)
      y_kmeans = model.fit_predict(satisfaction_evaluation_left)
```

```
[26]: # Convert the DataFrame to a NumPy array for correct indexing
      X = satisfaction_evaluation_left.values

      # Visualize the clusters
      plt.scatter(
```

```

    X[y_kmeans == 0, 0], X[y_kmeans == 0, 1],
    s=100, c='red', label='Cluster 1'
)
plt.scatter(
    X[y_kmeans == 1, 0], X[y_kmeans == 1, 1],
    s=100, c='blue', label='Cluster 2'
)
plt.scatter(
    X[y_kmeans == 2, 0], X[y_kmeans == 2, 1],
    s=100, c='green', label='Cluster 3'
)

# Add centroids
plt.scatter(
    model.cluster_centers[:, 0], model.cluster_centers[:, 1],
    s=300, c='yellow', marker='x', label='Centroids'
)

# Customize the plot
plt.title('Clusters of Employees Who Left')
plt.xlabel('Satisfaction Level')
plt.ylabel('Last Evaluation')
plt.legend()
plt.show()

```



#### 0.0.7 Clusters:

**Cluster 1 (Blue - High Satisfaction, High Evaluation):** Employees in this group are highly satisfied and performed well. Possible Reason for Leaving: Lack of growth opportunities or external offers despite strong performance.

**Cluster 2 (Red - Low Satisfaction, Medium Evaluation):** These employees have low satisfaction and moderate evaluations. Possible Reason for Leaving: Disengagement, dissatisfaction with work environment, or misaligned roles.

**Cluster 3 (Green - Low Satisfaction, High Evaluation):** Employees in this group have high evaluations but very low satisfaction. Possible Reason for Leaving: Burnout or lack of recognition for their efforts despite strong performance.

#### 0.0.8 4

Convert categorical variables to numerical variables

```
[27]: # 4 imbalanced data we have over here
class_distribution = df['left'].value_counts()
print(class_distribution)
```

```

left
0      10000
1       1991
Name: count, dtype: int64


```

```

[28]: # first we have to get the categorical data :
# Plot the distribution of the 'Class' column as a pie chart
plt.subplot(1, 2, 2)
class_distribution.plot(kind='pie', labels=['Stayed (0)', 'Left (1)'],
    ↪ colors=['blue', 'red'], autopct='%1.1f%%', startangle=90)
plt.title('Class Distribution')
plt.ylabel('')

# Show the plots
plt.tight_layout()
plt.show()

```



```

[29]: # Separate features and target
X = df.drop('left', axis=1) # Features
y = df['left']              # Target variable

# Identify categorical and numerical features
categorical_features = X.select_dtypes(include=['object']) # Categorical
    ↪ columns
numerical_features = X.select_dtypes(include=['number'])  # Numerical columns

```

```

# Apply one-hot encoding to categorical features
categorical_encoded = pd.get_dummies(categorical_features, drop_first=True).
    ↳ astype(int)

# Combine numerical and encoded categorical features
X_processed = pd.concat([numerical_features, categorical_encoded], axis=1)
X_processed

```

```

[29]:
      satisfaction_level  last_evaluation  number_project  \
0                0.38            0.53                2
1                0.80            0.86                5
2                0.11            0.88                7
3                0.72            0.87                5
4                0.37            0.52                2
...                ...                ...                ...
11995            0.90            0.55                3
11996            0.74            0.95                5
11997            0.85            0.54                3
11998            0.33            0.65                3
11999            0.50            0.73                4

      average_monthly_hours  time_spent_company  Work_accident  \
0                157                3                0
1                262                6                0
2                272                4                0
3                223                5                0
4                159                3                0
...                ...                ...                ...
11995            259                10                1
11996            266                10                0
11997            185                10                0
11998            172                10                0
11999            180                3                0

      promotion_last_5years  sales_RandD  sales_accounting  sales_hr  \
0                0                0                0                0
1                0                0                0                0
2                0                0                0                0
3                0                0                0                0
4                0                0                0                0
...                ...                ...                ...                ...
11995            1                0                0                0
11996            1                0                0                0
11997            1                0                0                0
11998            1                0                0                0
11999            0                0                0                0

```

	sales_management	sales_marketing	sales_product_mng	sales_sales	\
0	0	0	0	1	
1	0	0	0	1	
2	0	0	0	1	
3	0	0	0	1	
4	0	0	0	1	
...	...	...	...	...	
11995	1	0	0	0	
11996	1	0	0	0	
11997	1	0	0	0	
11998	0	1	0	0	
11999	0	0	0	0	

	sales_support	sales_technical	salary_low	salary_medium
0	0	0	1	0
1	0	0	0	1
2	0	0	0	1
3	0	0	1	0
4	0	0	1	0
...	...	...	...	...
11995	0	0	0	0
11996	0	0	0	0
11997	0	0	0	0
11998	0	0	0	0
11999	0	0	1	0

[11991 rows x 18 columns]

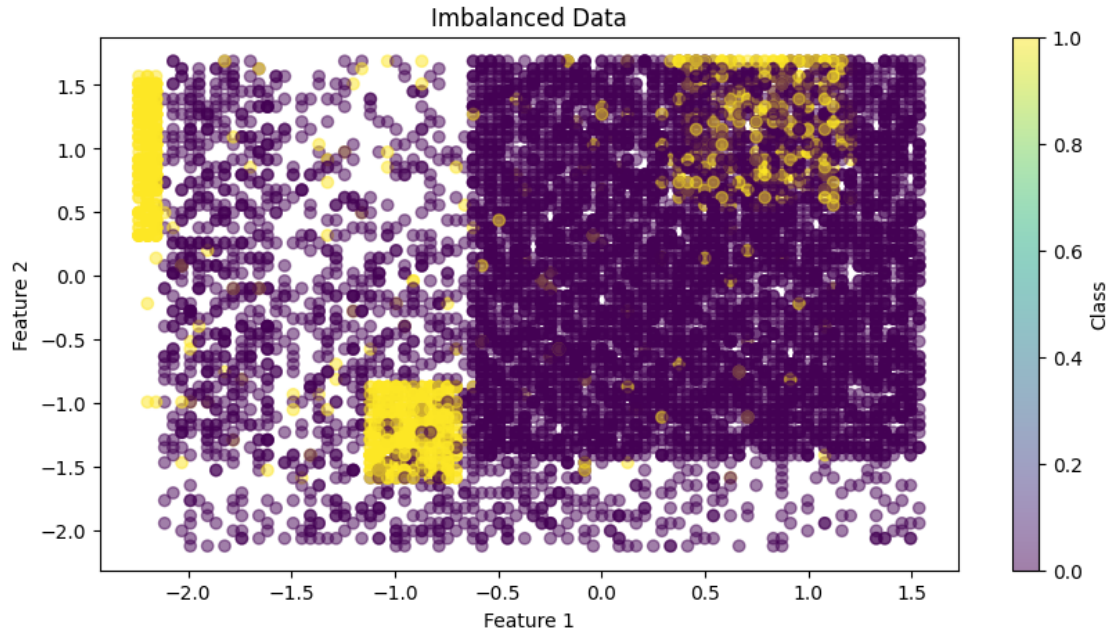
```
[90]: # Perform stratified split
X_train, X_test, y_train, y_test = train_test_split(
    X_processed, y,
    test_size=0.2,
    stratify=y,
    random_state=123
)

# Scale the numerical data
scaler = StandardScaler()
X_train_sc = scaler.fit_transform(X_train)
X_test_sc = scaler.transform(X_test)
```

```
[91]: # Scatter plot for the imbalanced data
plt.figure(figsize=(10, 5))
plt.scatter(X_train_sc[:, 0], X_train_sc[:, 1], c=y_train, alpha=0.5,
            cmap='viridis', marker='o')
plt.title('Imbalanced Data')
plt.xlabel('Feature 1')
```

```
plt.ylabel('Feature 2')
plt.colorbar(label='Class')
```

[91]: <matplotlib.colorbar.Colorbar at 0x31b697c70>



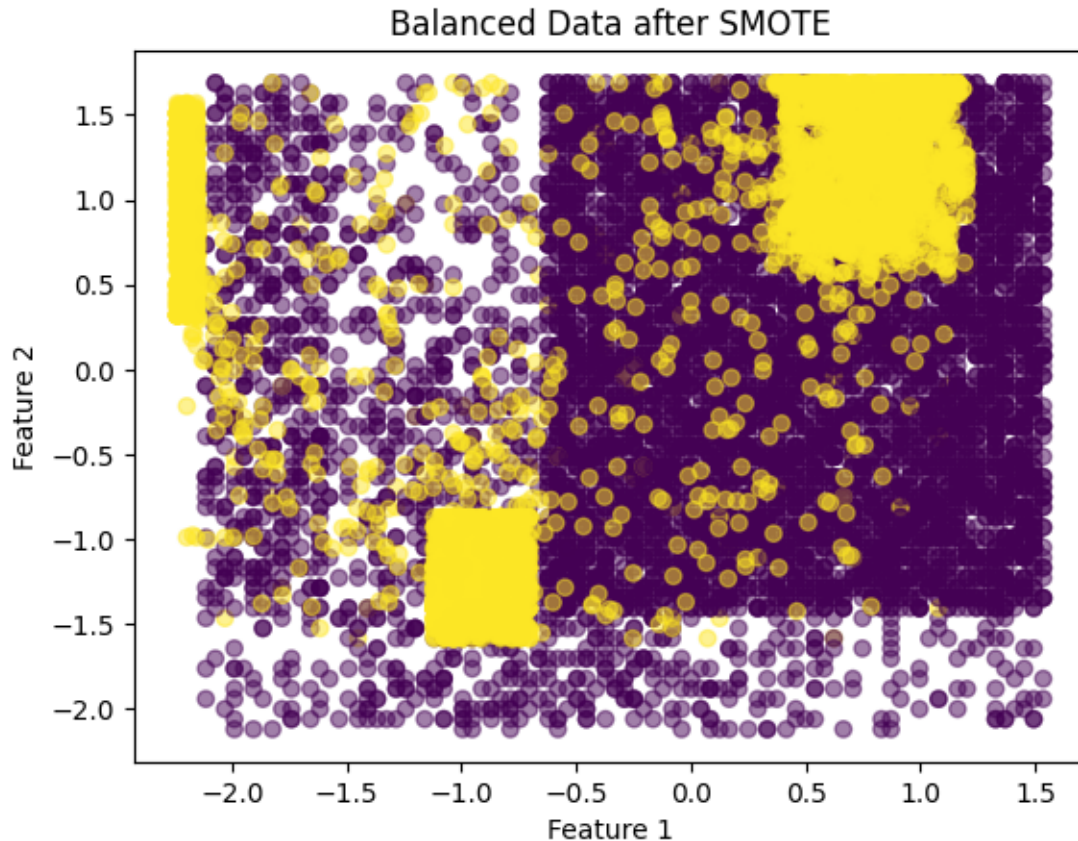
[92]: *# now I have them combined , so I can proceed with the rest of the tasks*

```
# Apply SMOTE
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train_sc, y_train)
```

```
plt.scatter(X_train_smote[:, 0], X_train_smote[:, 1], c=y_train_smote, alpha=0.
↪5, cmap="viridis", marker='o')
plt.title('Balanced Data after SMOTE')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
```

[94]: Text(0, 0.5, 'Feature 2')





### 0.0.9 5

Perform 5-fold cross-validation model training and evaluate performance

### 0.0.10 Logistic Regression Model

```
[109]: # Initialize the Logistic Regression model
logistic_model = LogisticRegression(random_state=123)

# Perform 5-fold cross-validation and predict probabilities for the positive
# class
y_prob_cv = cross_val_predict(logistic_model, X_train_smote, y_train_smote,
                              cv=5, method='predict_proba')[:, 1]

# Compute ROC curve
fpr, tpr, thresholds = roc_curve(y_train_smote, y_prob_cv)

# Compute AUC
roc_auc = auc(fpr, tpr)
```

```

# Compute Youden's J statistic and the optimal threshold
youden_j = tpr - fpr
optimal_threshold_index = np.argmax(youden_j)
optimal_threshold = thresholds[optimal_threshold_index]

# Print classification report
y_pred_cv = (y_prob_cv >= optimal_threshold).astype(int) # Use optimal_
↳threshold for predictions
print(f"Logistic Regression Optimal Threshold: {optimal_threshold:.4f}")
print("Logistic Regression Classification Report:")
print(classification_report(y_train_smote, y_pred_cv))

# Generate the classification report as a dictionary
report = classification_report(y_train_smote, y_pred_cv, output_dict=True)

# Convert the dictionary into a pandas DataFrame
report_df = pd.DataFrame(report).transpose()

# Filter for classes 0 and 1 (employees on the left class)
filtered_report = report_df.loc[["0", "1"], ["precision", "recall", "f1-score"]]

# Compute the confusion matrix
cm = confusion_matrix(y_train_smote, y_pred_cv)

# Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Class 0',
↳'Class 1'])
disp.plot(cmap="summer", values_format="d") # Using 'summer' colormap as
↳requested
plt.title("Confusion Matrix - Logistic Regression")
plt.ylabel("True Label")
plt.xlabel("Predicted Label")
plt.show()


# Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.scatter(fpr[optimal_threshold_index], tpr[optimal_threshold_index],
↳color='red', label=f'Optimal Threshold = {optimal_threshold:.4f}')
plt.plot([0, 1], [0, 1], color='grey', linestyle='--', lw=2) # Diagonal line
↳for random guessing
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Logistic Regression')
plt.legend(loc='lower right')
plt.grid(alpha=0.3)
plt.show()


```

Logistic Regression Optimal Threshold: 0.4524

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.87	0.73	0.79	7999
1	0.77	0.89	0.82	7999
accuracy			0.81	15998
macro avg	0.82	0.81	0.81	15998
weighted avg	0.82	0.81	0.81	15998





### 0.0.11 Random Forest Classifier Model

```
[122]: # Initialize the Random Forest model
rf_model = RandomForestClassifier(random_state=123)

# Perform 5-fold cross-validation and predict probabilities for the positive_
↪class
y_prob_cv_rf = cross_val_predict(rf_model, X_train_smote, y_train_smote, cv=5, ↪
↪method='predict_proba')[:, 1]

# Compute ROC curve
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_train_smote, y_prob_cv_rf)

# Compute AUC
roc_auc_rf = auc(fpr_rf, tpr_rf)

# Compute Youden's J statistic and the optimal threshold
youden_j_rf = tpr_rf - fpr_rf
optimal_threshold_index_rf = np.argmax(youden_j_rf)
```

```

optimal_threshold_rf = thresholds_rf[optimal_threshold_index_rf]

# Generate predictions based on the optimal threshold
y_pred_cv_rf = (y_prob_cv_rf >= optimal_threshold_rf).astype(int)

# Print classification report
print(f"Random Forest Optimal Threshold: {optimal_threshold_rf:.4f}")
print("Random Forest Classification Report:")
print(classification_report(y_train_smote, y_pred_cv_rf))

# Generate the classification report as a dictionary
report_rf = classification_report(y_train_smote, y_pred_cv_rf, output_dict=True)

# Convert the dictionary into a pandas DataFrame
report_df_rf = pd.DataFrame(report_rf).transpose()

# Filter for classes 0 and 1
filtered_report_rf = report_df_rf.loc[["0", "1"], ["precision", "recall", "f1-score"]]

# Compute the confusion matrix
cm_rf = confusion_matrix(y_train_smote, y_pred_cv_rf)

# Display the confusion matrix
disp_rf = ConfusionMatrixDisplay(confusion_matrix=cm_rf, display_labels=['Class 0', 'Class 1'])
disp_rf.plot(cmap="summer", values_format="d") # Using 'summer' colormap
plt.title("Confusion Matrix - Random Forest")
plt.ylabel("True Label")
plt.xlabel("Predicted Label")
plt.show()


# Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, color='green', lw=2, label=f'ROC Curve (AUC = {roc_auc_rf:.2f})')
plt.scatter(fpr_rf[optimal_threshold_index_rf], tpr_rf[optimal_threshold_index_rf], color='red', label=f'Optimal Threshold = {optimal_threshold_rf:.4f}')
plt.plot([0, 1], [0, 1], color='grey', linestyle='--', lw=2) # Diagonal line for random guessing
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Random Forest')
plt.legend(loc='lower right')
plt.grid(alpha=0.3)
plt.show()

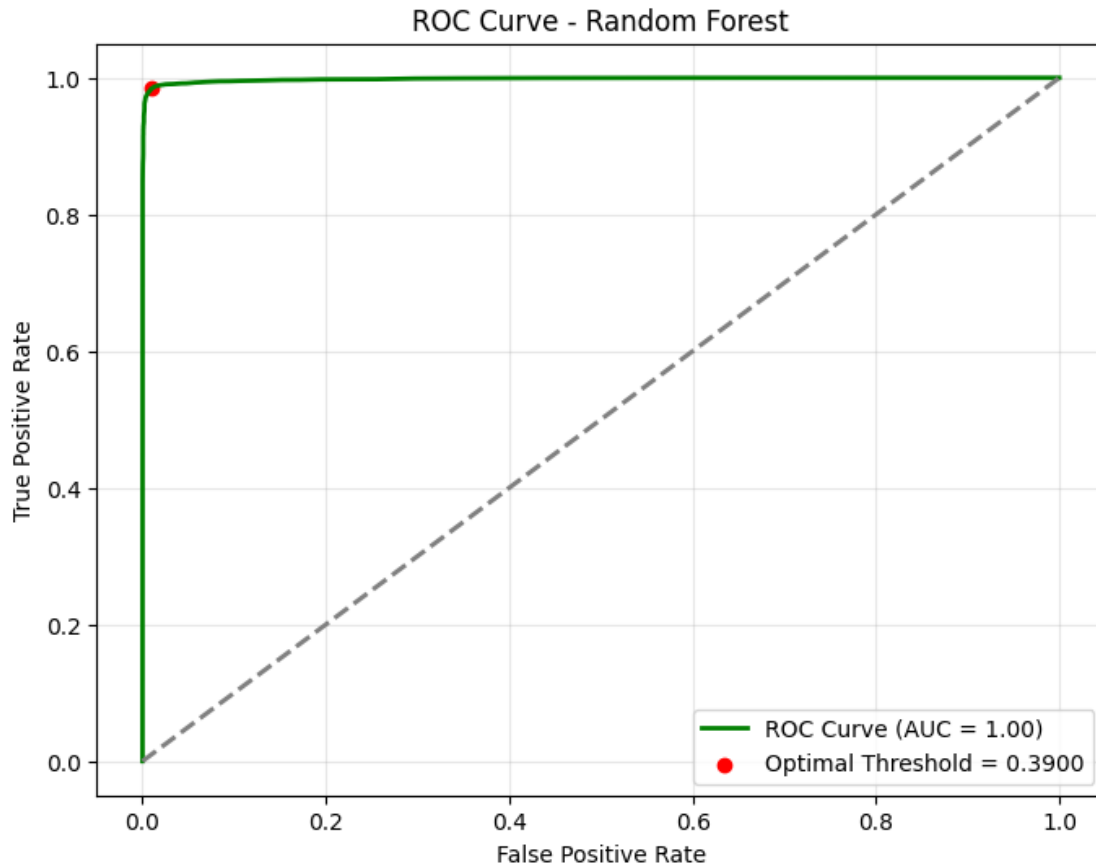
```

Random Forest Optimal Threshold: 0.3900

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	7999
1	0.99	0.98	0.99	7999
accuracy			0.99	15998
macro avg	0.99	0.99	0.99	15998
weighted avg	0.99	0.99	0.99	15998





### 0.0.12 Gradient Boosting Classifier Model

```
[123]: # Initialize the Gradient Boosting model
gb_model = GradientBoostingClassifier(random_state=123)

# Perform 5-fold cross-validation and predict probabilities for the positive
# class
y_prob_cv_gb = cross_val_predict(gb_model, X_train_smote, y_train_smote, cv=5,
# method='predict_proba')[:, 1]

# Compute ROC curve
fpr_gb, tpr_gb, thresholds_gb = roc_curve(y_train_smote, y_prob_cv_gb)

# Compute AUC
roc_auc_gb = auc(fpr_gb, tpr_gb)

# Compute Youden's J statistic and the optimal threshold
youden_j_gb = tpr_gb - fpr_gb
optimal_threshold_index_gb = np.argmax(youden_j_gb)
```

```

optimal_threshold_gb = thresholds_gb[optimal_threshold_index_gb]

# Generate predictions based on the optimal threshold
y_pred_cv_gb = (y_prob_cv_gb >= optimal_threshold_gb).astype(int)

# Print classification report
print(f"Gradient Boosting Optimal Threshold: {optimal_threshold_gb:.4f}")
print("Gradient Boosting Classification Report:")
print(classification_report(y_train_smote, y_pred_cv_gb))

# Generate the classification report as a dictionary
report_gb = classification_report(y_train_smote, y_pred_cv_gb, output_dict=True)

# Convert the dictionary into a pandas DataFrame
report_df_gb = pd.DataFrame(report_gb).transpose()

# Filter for classes 0 and 1
filtered_report_gb = report_df_gb.loc[["0", "1"], ["precision", "recall",
↪ "f1-score"]]

# Compute the confusion matrix
cm_gb = confusion_matrix(y_train_smote, y_pred_cv_gb)

# Display the confusion matrix
disp_gb = ConfusionMatrixDisplay(confusion_matrix=cm_gb, display_labels=['Class_
↪ 0', 'Class 1'])
disp_gb.plot(cmap="plasma", values_format="d") # Using 'plasma' colormap for a
↪ vibrant look
plt.title("Confusion Matrix - Gradient Boosting")
plt.ylabel("True Label")
plt.xlabel("Predicted Label")
plt.show()

# Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_gb, tpr_gb, color='purple', lw=2, label=f'ROC Curve (AUC =
↪ {roc_auc_gb:.2f})')
plt.scatter(fpr_gb[optimal_threshold_index_gb],
↪ tpr_gb[optimal_threshold_index_gb], color='red', label=f'Optimal Threshold =
↪ {optimal_threshold_gb:.4f}')
plt.plot([0, 1], [0, 1], color='grey', linestyle='--', lw=2) # Diagonal line
↪ for random guessing
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Gradient Boosting Classifier')
plt.legend(loc='lower right')

```




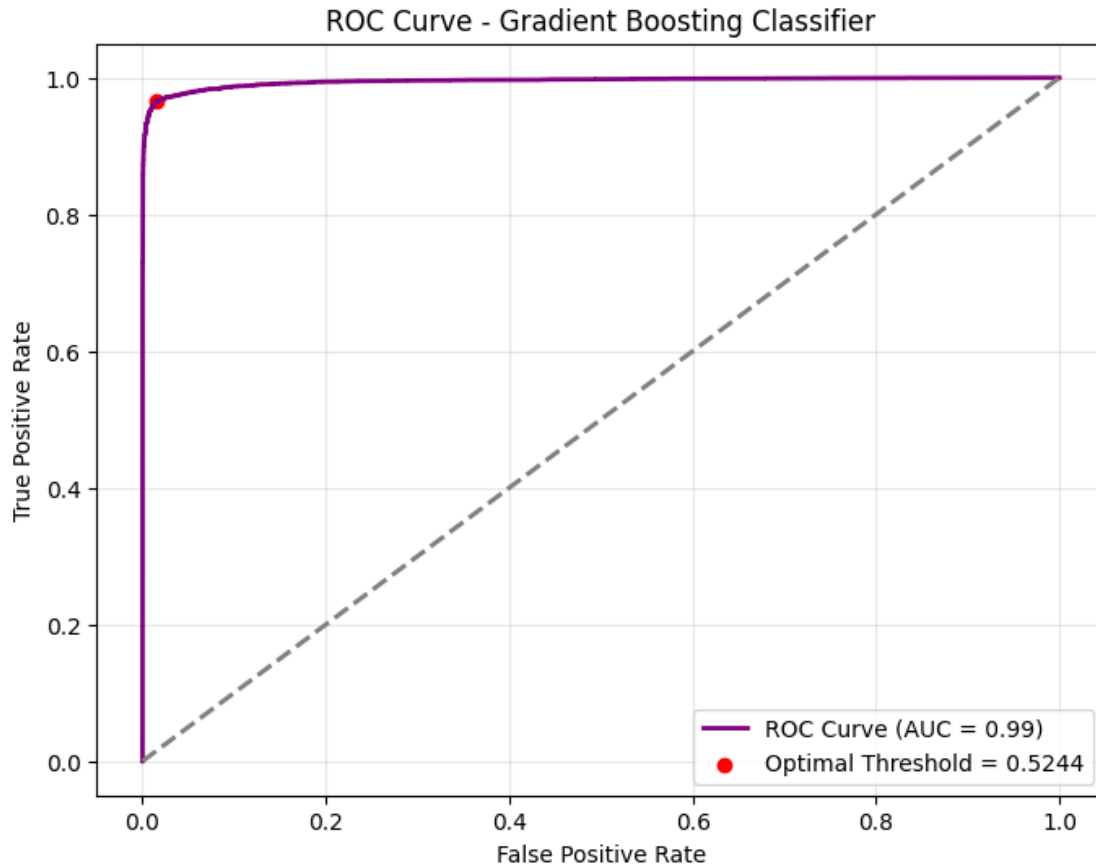
```
plt.grid(alpha=0.3)
plt.show()
```

Gradient Boosting Optimal Threshold: 0.5244

Gradient Boosting Classification Report:

	precision	recall	f1-score	support
0	0.97	0.98	0.98	7999
1	0.98	0.97	0.98	7999
accuracy			0.98	15998
macro avg	0.98	0.98	0.98	15998
weighted avg	0.98	0.98	0.98	15998





**Best Model: Gradient Boosting.**

Offers high AUC (0.99) with balanced Precision and Recall, making it reliable and less prone to overfitting.

**Metric to Optimize: Recall** (to capture as many employees who left as possible).

```
[114]: # Fit the Gradient Boosting model on the entire training data
gb_model.fit(X_train_smote, y_train_smote)

# Predict probabilities for the test data
y_prob_test = gb_model.predict_proba(X_test)[: , 1]

# Add the probabilities to a DataFrame for easy categorization
test_results = pd.DataFrame({
    'EmployeeID': X_test.index,
    'Turnover_Probability': y_prob_test
})
```

```

# Define a function to categorize employees into risk zones
def categorize_risk(prob):
    if prob < 0.2:
        return 'Safe Zone (Green)'
    elif 0.2 <= prob < 0.6:
        return 'Low-Risk Zone (Yellow)'
    elif 0.6 <= prob < 0.9:
        return 'Medium-Risk Zone (Orange)'
    else:
        return 'High-Risk Zone (Red)'

# Categorize employees into risk zones
test_results['Risk_Zone'] = test_results['Turnover_Probability'].
    ↪ apply(categorize_risk)

# Ensure Risk_Zone column has a specific order
risk_zone_order = ['Safe Zone (Green)', 'Low-Risk Zone (Yellow)', 'Medium-Risk_
    ↪ Zone (Orange)', 'High-Risk Zone (Red)']
test_results['Risk_Zone'] = pd.Categorical(test_results['Risk_Zone'],
    ↪ categories=risk_zone_order, ordered=True)

# Display the first few rows of the categorized results
print(test_results.head())

plt.figure(figsize=(8, 6))
sns.countplot(
    x='Risk_Zone',
    data=test_results,
    palette={
        'Safe Zone (Green)': 'green',
        'Low-Risk Zone (Yellow)': 'yellow',
        'Medium-Risk Zone (Orange)': 'orange',
        'High-Risk Zone (Red)': 'red'
    }
)
plt.title('Risk Zone Distribution of Employees', fontsize=16)
plt.xlabel('Risk Zone', fontsize=14)
plt.ylabel('Number of Employees', fontsize=14)
plt.xticks(rotation=15)
plt.grid(alpha=0.3)
plt.show()

```

	EmployeeID	Turnover_Probability	Risk_Zone
0	8578	0.802393	Medium-Risk Zone (Orange)
1	5756	0.674578	Medium-Risk Zone (Orange)
2	3994	0.974625	High-Risk Zone (Red)

```

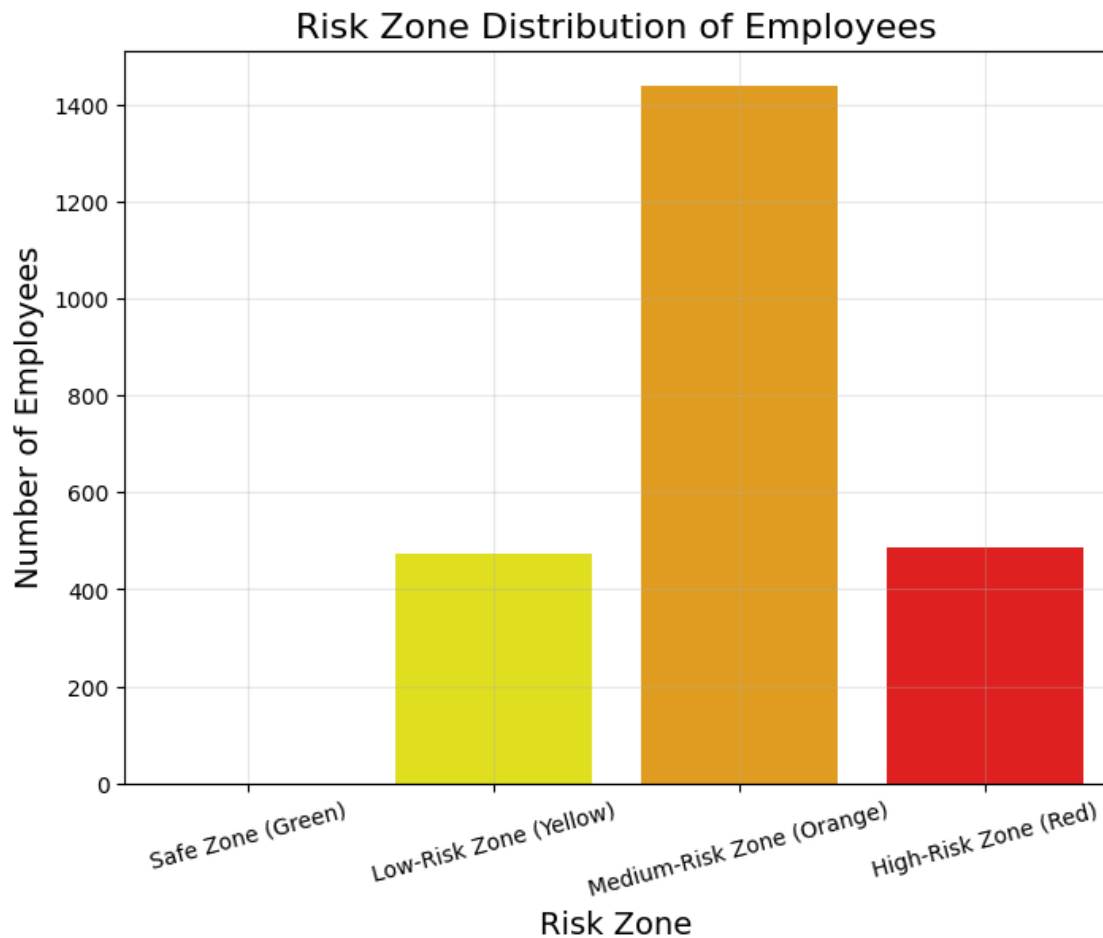
3         1784                0.417309    Low-Risk Zone (Yellow)
4         10508               0.846878    Medium-Risk Zone (Orange)

/opt/anaconda3/envs/simplienv/lib/python3.9/site-packages/sklearn/base.py:486:
UserWarning: X has feature names, but GradientBoostingClassifier was fitted
without feature names
  warnings.warn(
/var/folders/y3/g8cf6rzj7wz97zptmgzfqqffc0000gn/T/ipykernel_14183/3539641596.py:3
7: FutureWarning:

```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(
```



### 0.0.13 Retention Strategies:

1. **Safe Zone (Green) (<20%)** Maintain satisfaction with rewards, recognition, and growth opportunities. Regular check-ins to address potential concerns.
2. **Low-Risk Zone (Yellow) (20%-60%)** Boost engagement with involvement in decisions and occasional perks. Address concerns through surveys or one-on-ones.
3. **Medium-Risk Zone (Orange) (60%-90%)** Provide growth paths, better compensation, and training. Act on employee feedback to resolve dissatisfaction.
4. **High-Risk Zone (Red) (>90%)** Address critical issues (pay, workload) immediately. Create personalized plans and fix root causes quickly.

[ ]: